

CS341: Computer Architecture Lab

Assignment 4

Report

Sudhir Kumar (170050053)



Department of Computer Science and Engineering
Indian Institute of Technology Bombay
2021-2022

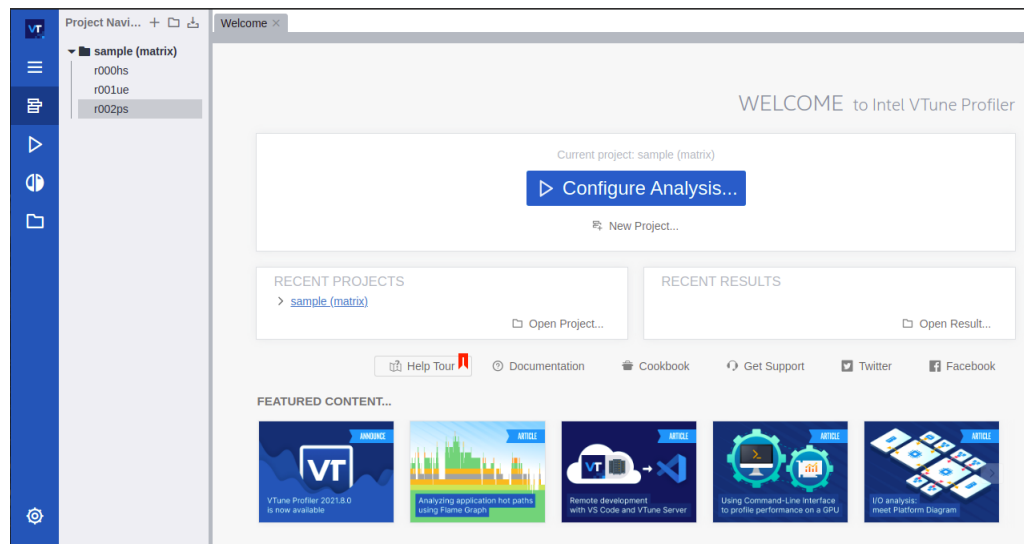
Contents

1	Part 0: Getting Things Ready	1
1.1	Intel VTune Profiler Installation	1
1.2	Challenges during installation	1
1.3	Docker installation	1
2	Part 1: Profiling with VTune	3
2.1	Performance Snapshot	3
2.1.1	bfs.cpp	3
2.1.2	matrix_multi.cpp	5
2.1.3	matrix_multi_2.cpp	7
2.1.4	quicksort.cpp	9
2.2	Hotspots	11
2.2.1	bfs.cpp	11
2.2.2	matrix_multi.cpp	13
2.2.3	matrix_multi_2.cpp	15
2.2.4	quicksort.cpp	17
	Bibliography	18
3	Part 2: Simulating with ChampSim	19
3.1	Preparing the traces	19
3.2	baseline	19
3.3	direct-mapped/: Effect of using Direct-Mapped Cache at all levels	20
3.4	fully-associative/: Effect of using Fully-Associative Cache at all levels	20
3.5	reduced-size/: Effect of halving the size of the caches at all levels	21
3.6	doubled-size/: Effect of doubling the size of the caches at all levels	22
3.7	doubled-mshr/: Effect of doubling the number of the MSHRs at all levels	22
3.8	reduced-mshr/:Effect of halving the number of MSHRs at all levels	23

1. Part 0: Getting Things Ready

1.1 Intel VTune Profiler Installation

followed the guidance according to given link and installed Intel VTune Profiler on ubuntu 20.04.



1.2 Challenges during installation

It was not much difficult to install. I just followed all the instructions given at the site provided in problem statements and relative websites. After I downloaded the package followed the instructions of installer. It took me almost 1 hour, 30 minutes to understand, download and install the application. But again it took almost one and half hours to check if I have downloaded correctly and how to open the application. So total of 3 hours for installation and to get it running.

1.3 Docker installation

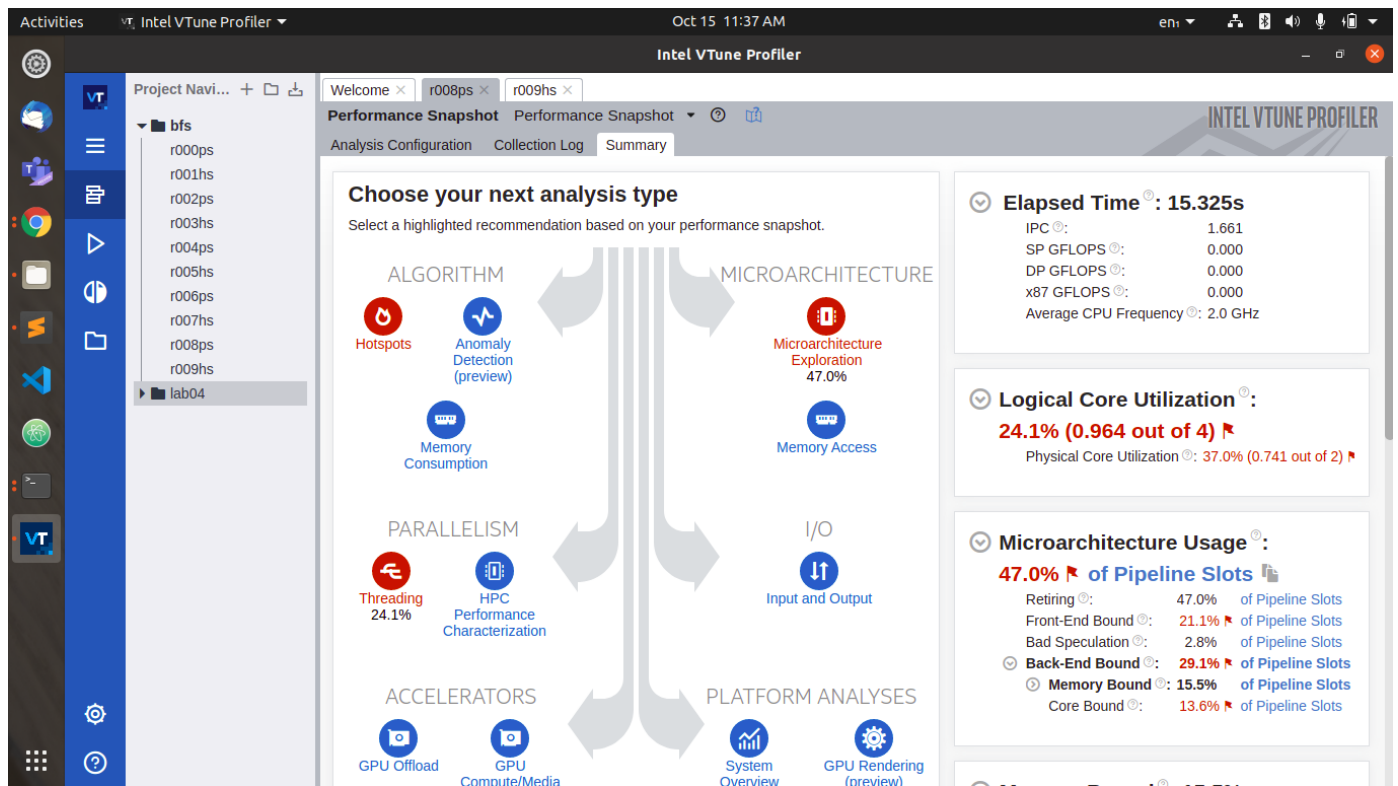
Installed using guide from internet and pulled 0xd3ba/champsim-lab

2. Part 1: Profiling with VTune

located provided programs and run with -g -O2 flags to create executable after suitable modification in codes to get execution time as per guided.

2.1 Performance Snapshot

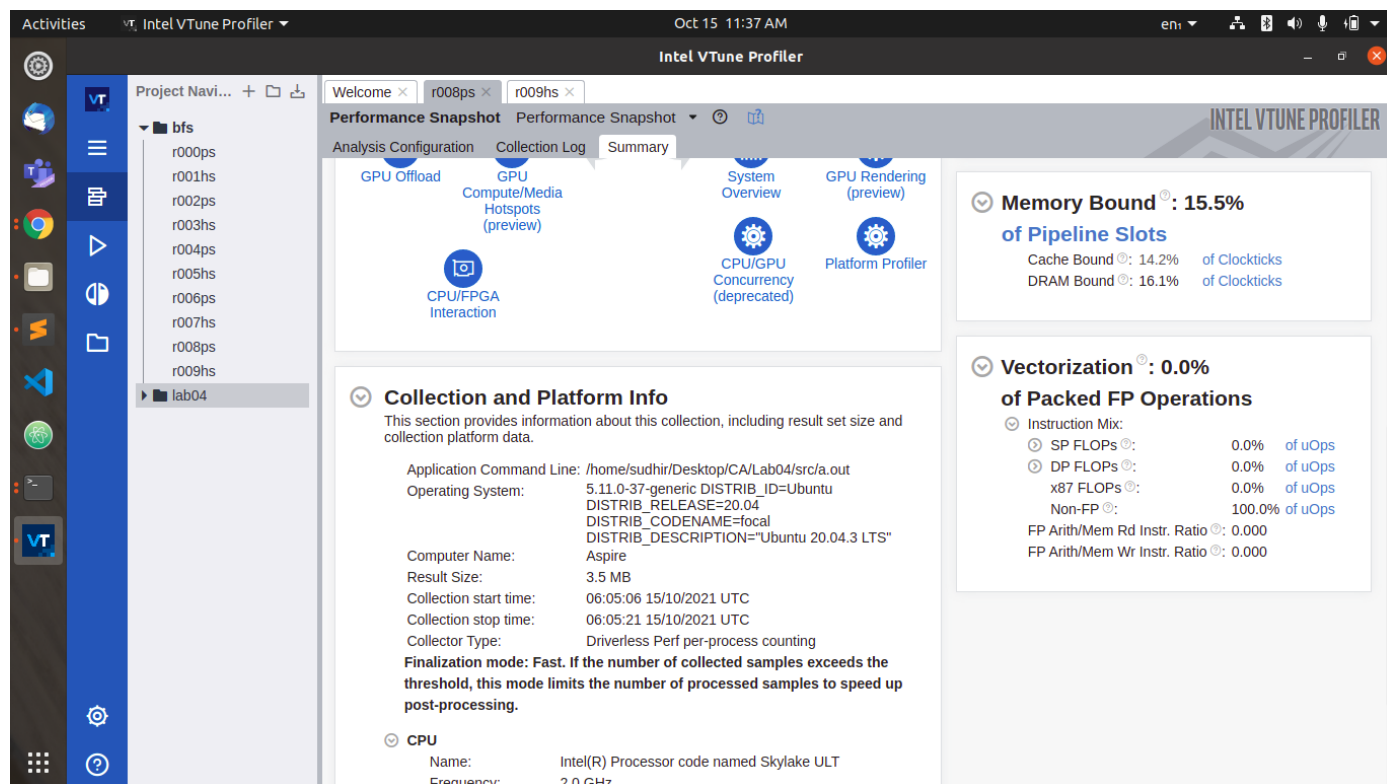
2.1.1 bfs.cpp



IPC = 1.661

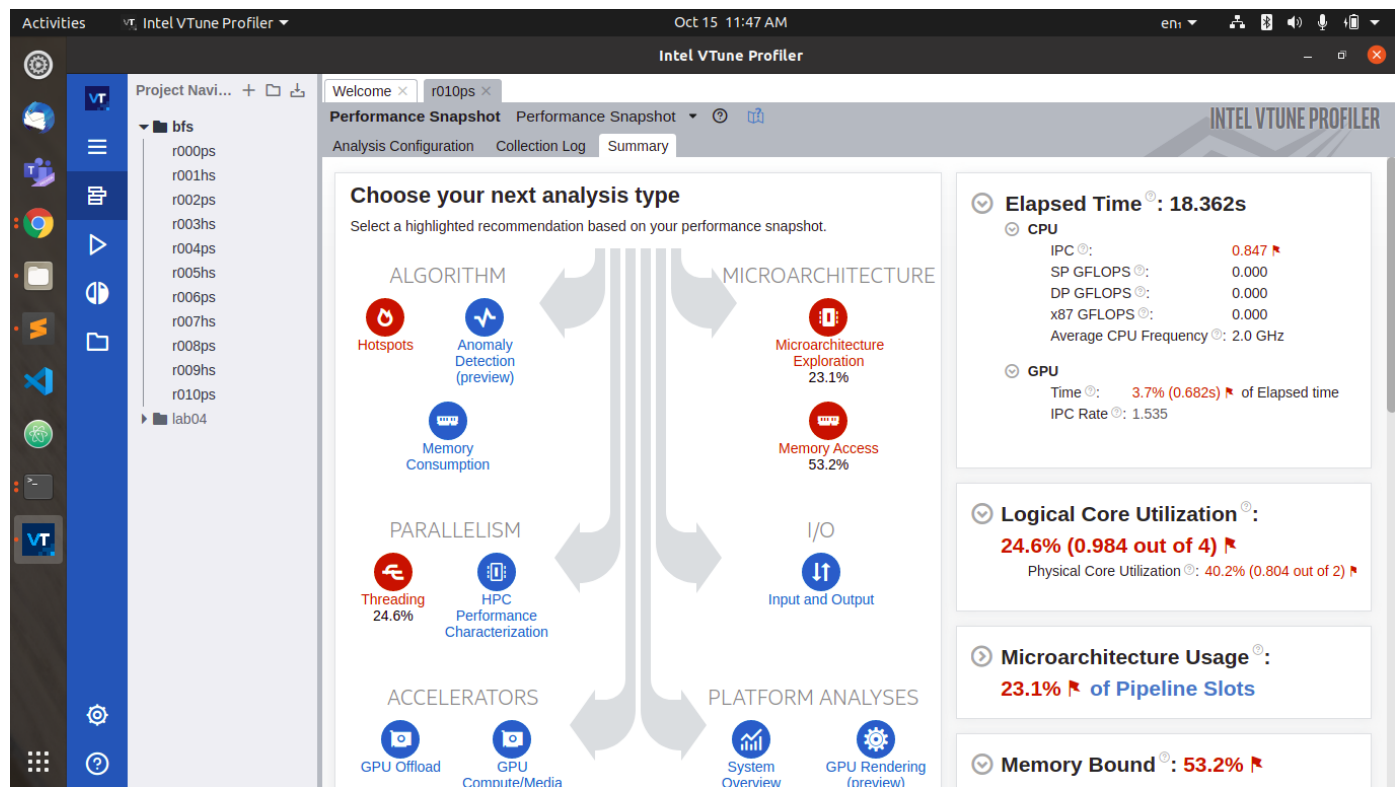
Logical core utilization = 24.1% (0.964 out of 4)

Physical core utilization = 37.0% (0.741 out of 2)



Memory Bound = 15.5 %

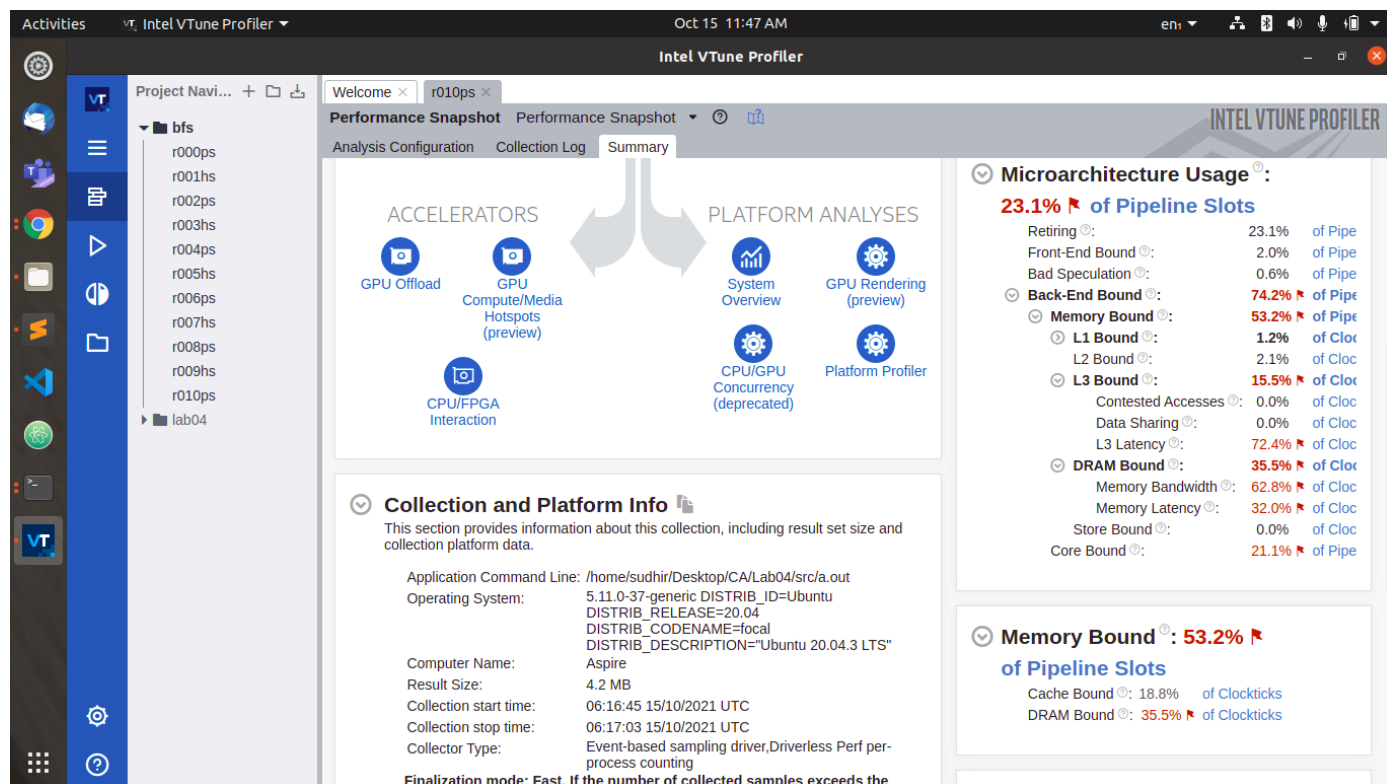
2.1.2 matrix_multi.cpp



IPC = 0.847

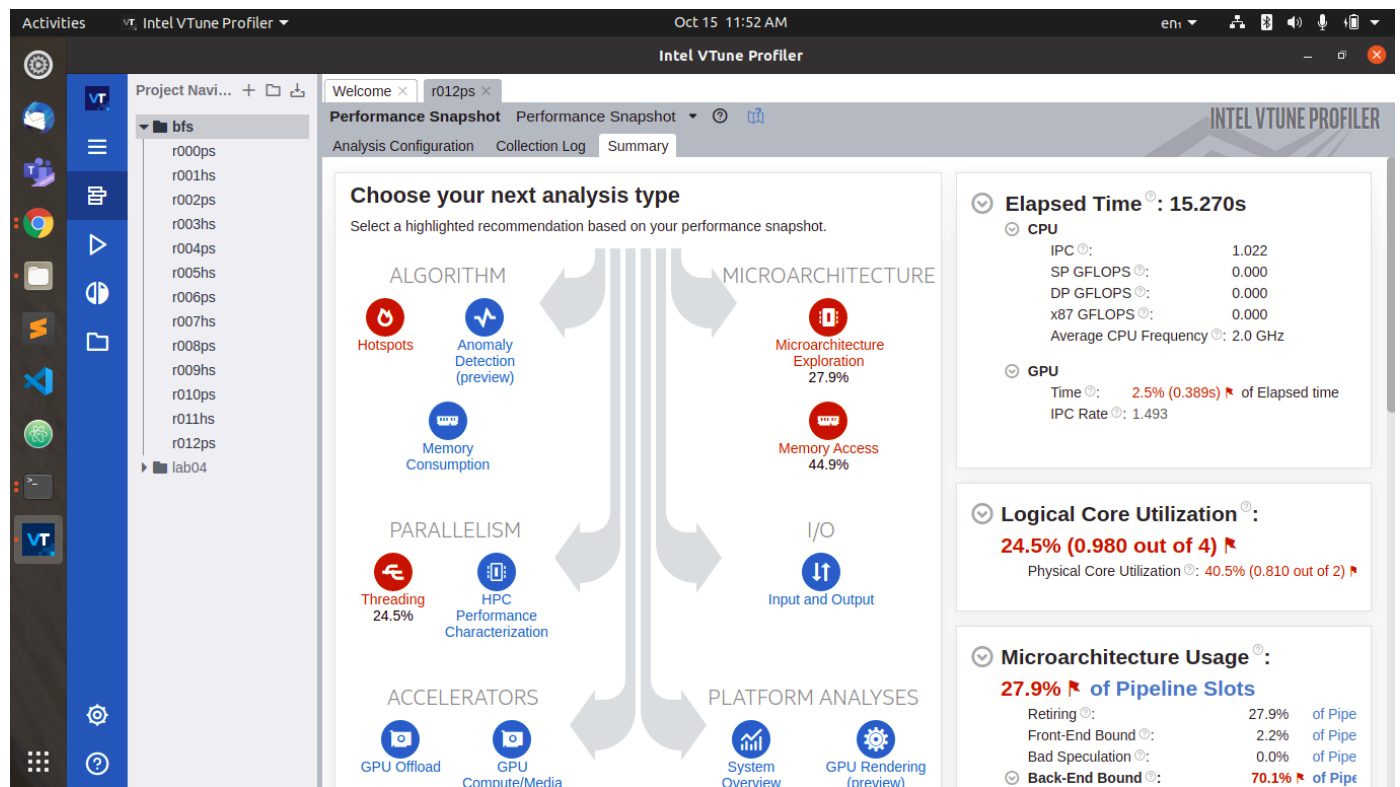
Logical core utilization = 24.6% (0.984 out of 4)

Physical core utilization = 40.2% (0.804 out of 2)



Memory Bound = 53.2 %

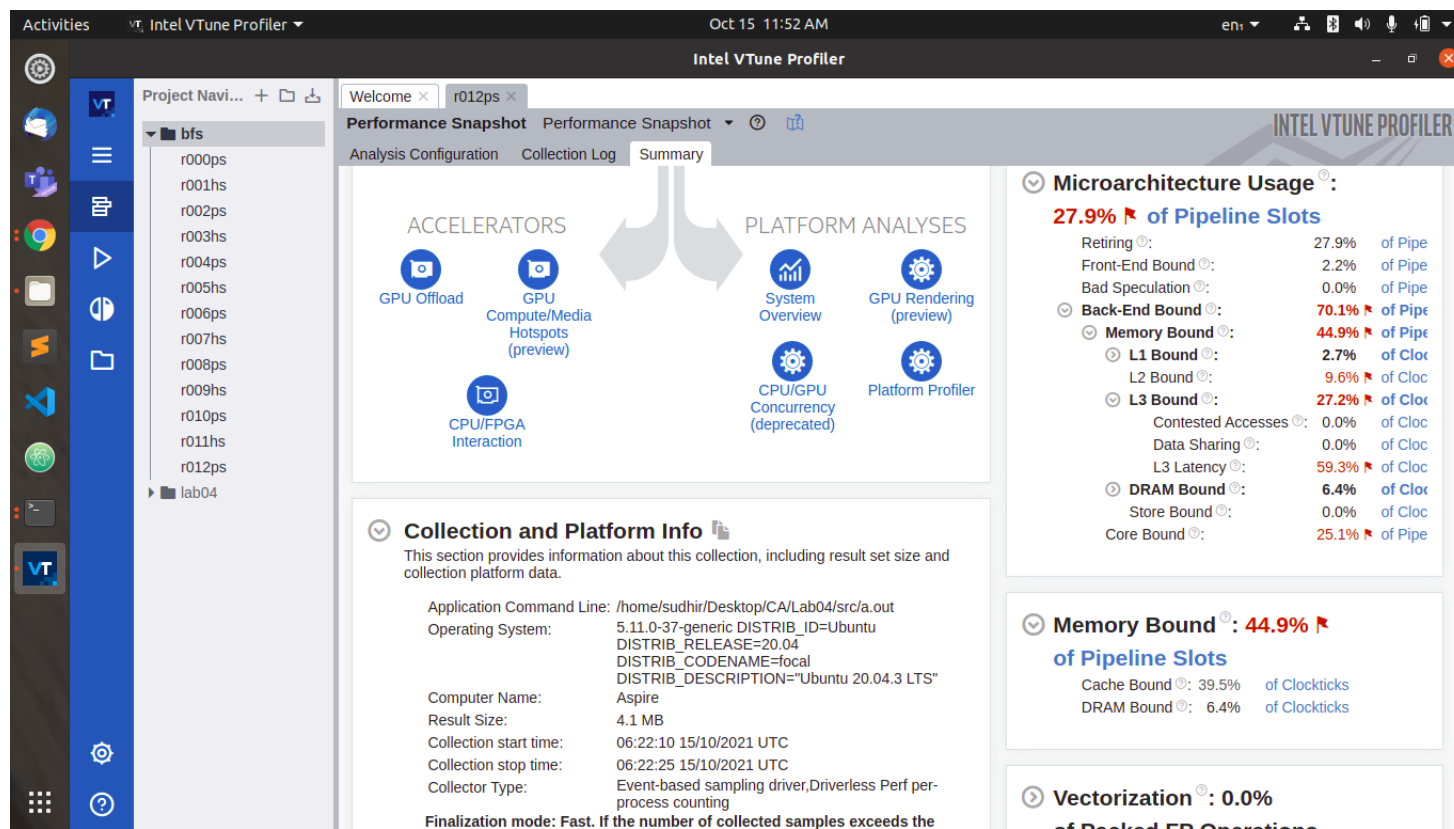
2.1.3 matrix_multi_2.cpp



IPC = 1.022

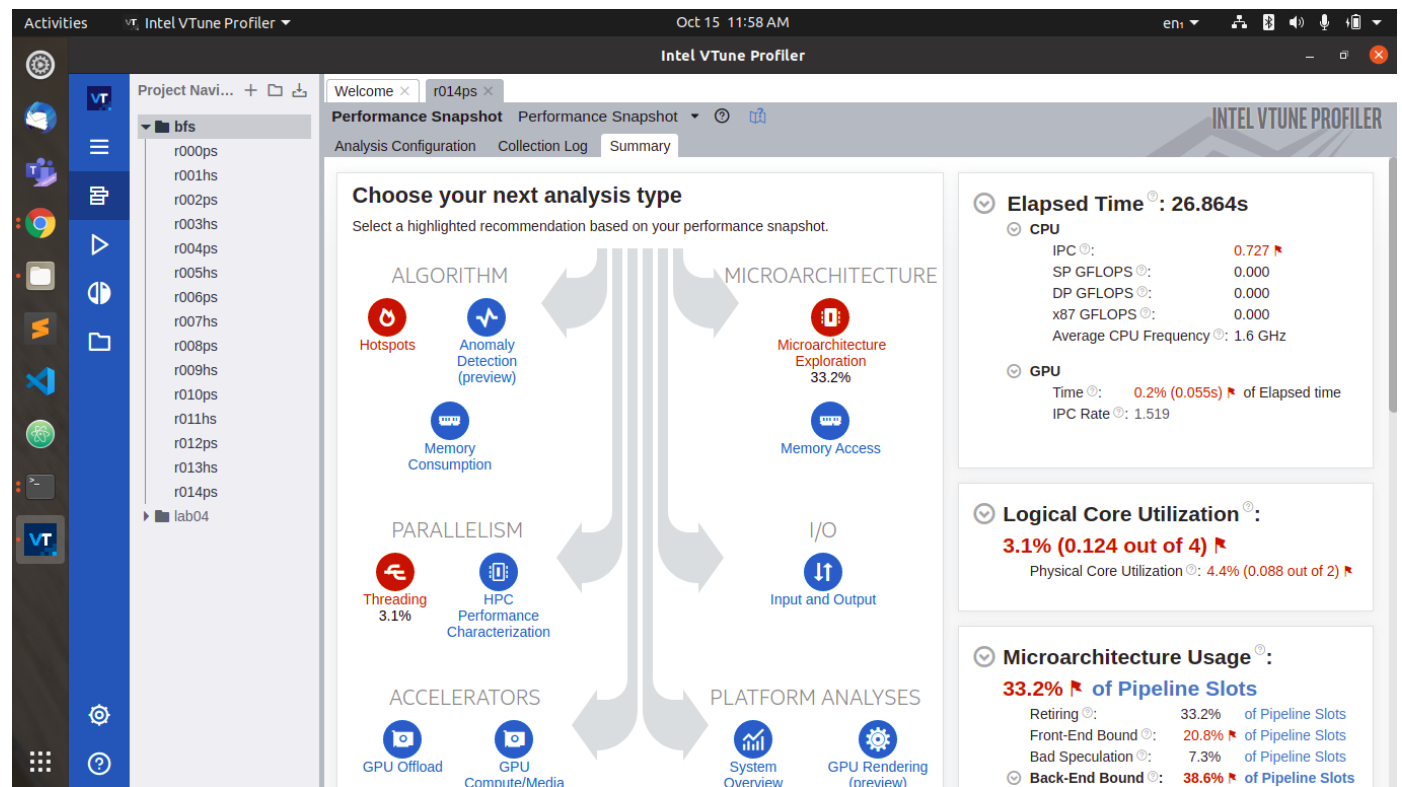
Logical core utilization = 24.5% (0.980 out of 4)

Physical core utilization = 40.5% (0.810 out of 2)



Memory Bound = 44.9 %

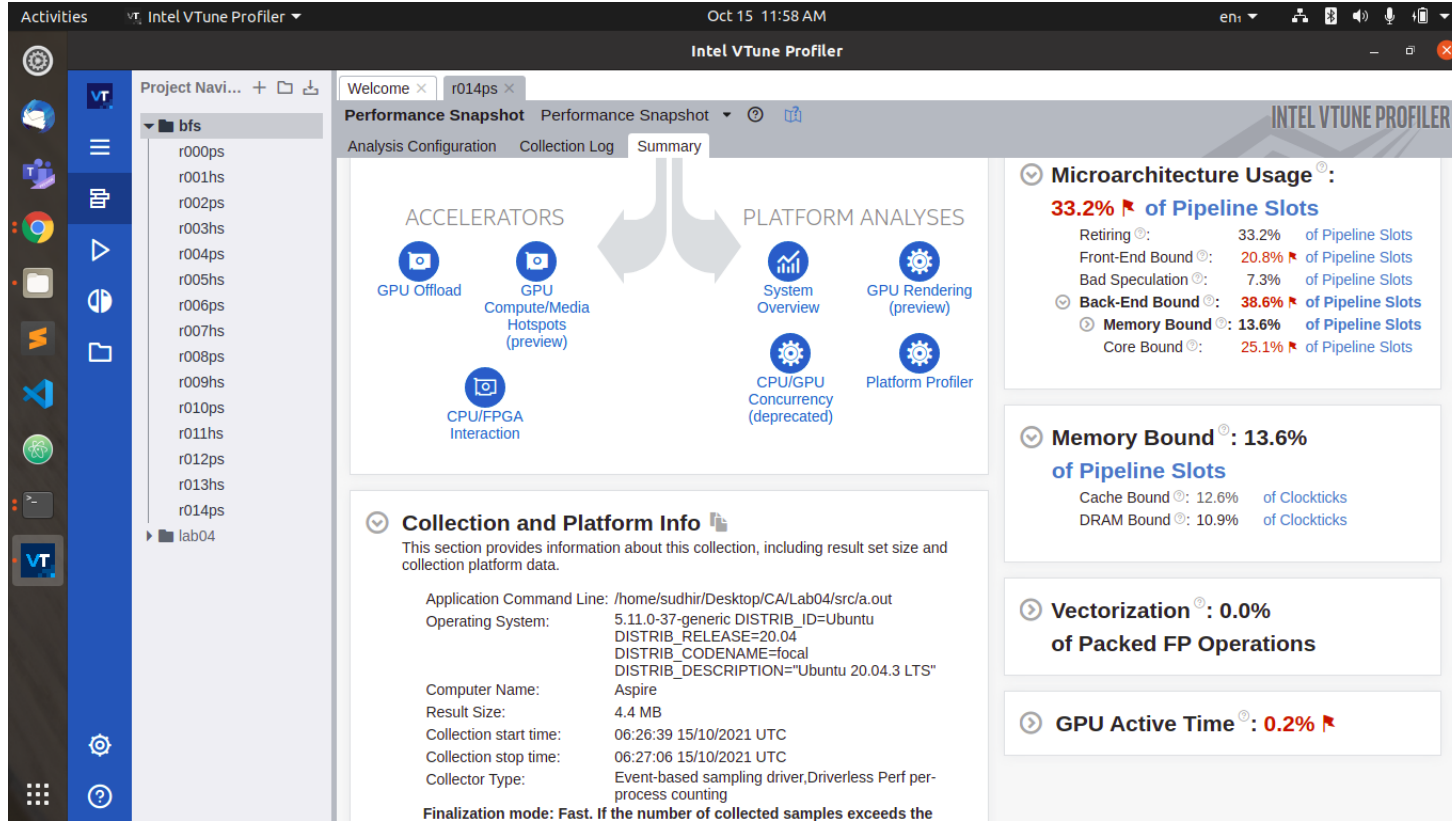
2.1.4 quicksort.cpp



IPC = 0.727

Logical core utilization = 3.1% (0.124 out of 4)

Physical core utilization = 4.4% (0.088 out of 2)



Memory Bound = 13.6 %

2.2 Hotspots

2.2.1 bfs.cpp

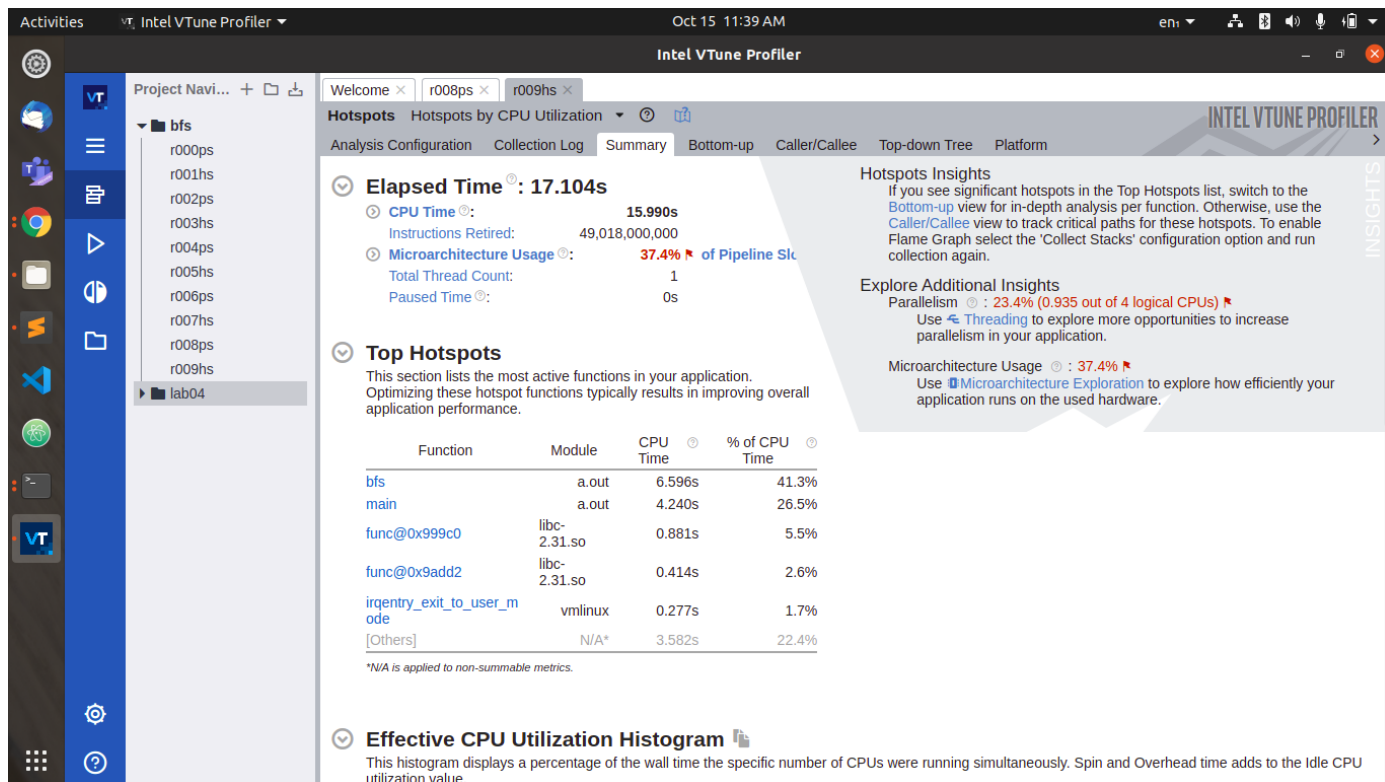


Table 2.1: Top Hotspots

Function	Module	CPU Time	% of CPU Time
bfs	a.out	6.596s	41.3%
main	a.out	4.240s	26.5%
func@0x999c0	libc-2.31.so	0.881s	5.5%
func@0x9add2	libc-2.31.so	0.414s	2.6%
irqentry_exit_to_user_mode	vmlinux	0.277s	1.7%
[Others]	N/A*	3.582s	22.4%

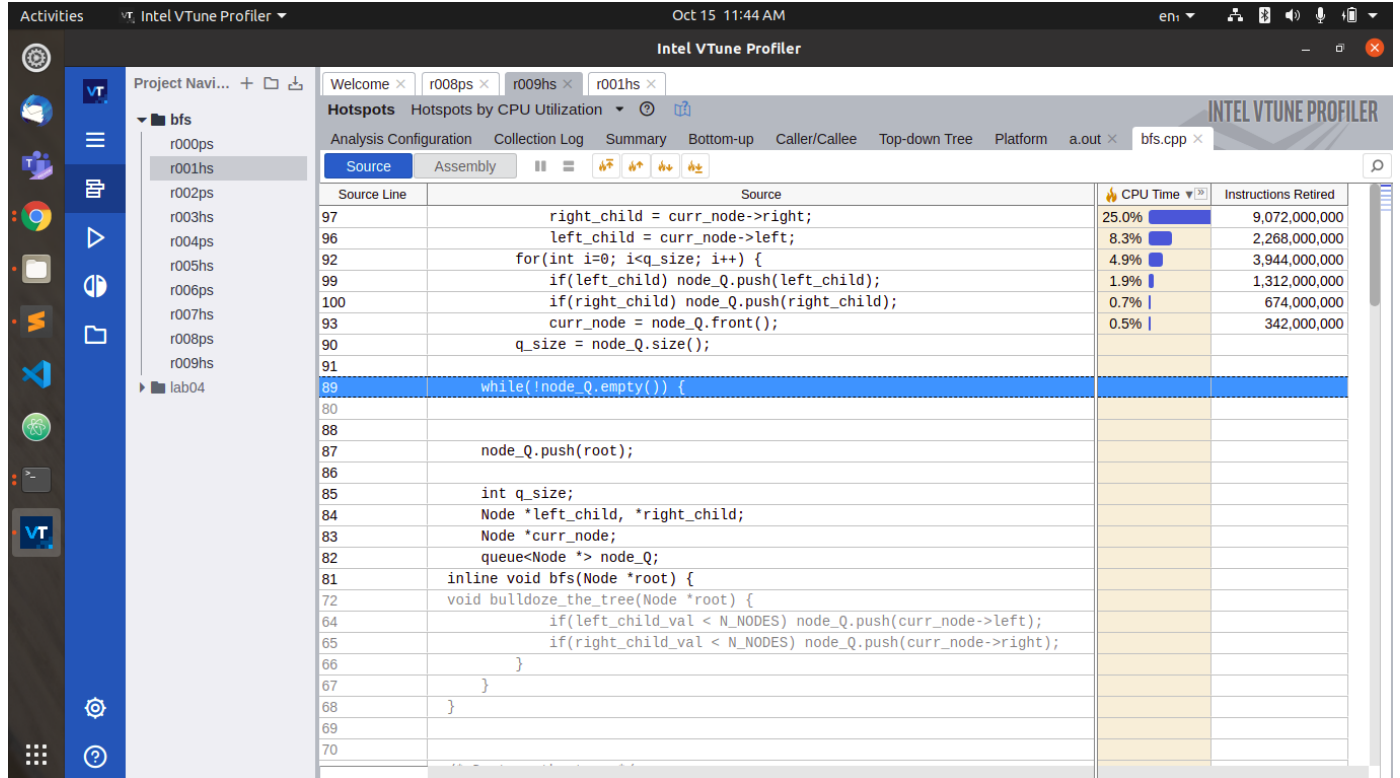


Table 2.2: Statements and % CPU Time

Source	% CPU Time
right_child = curr_node->right;	25.0%
left_child = curr_node->left;	8.3%
for(int i=0; i<q_size; i++) {	4.9%
if(left_child) node_Q.push(left_child);	1.9%
if(right_child) node_Q.push(right_child);	0.7%
curr_node = node_Q.front();	0.5%
bfs(root);	26.5%

2.2.2 matrix_multi.cpp

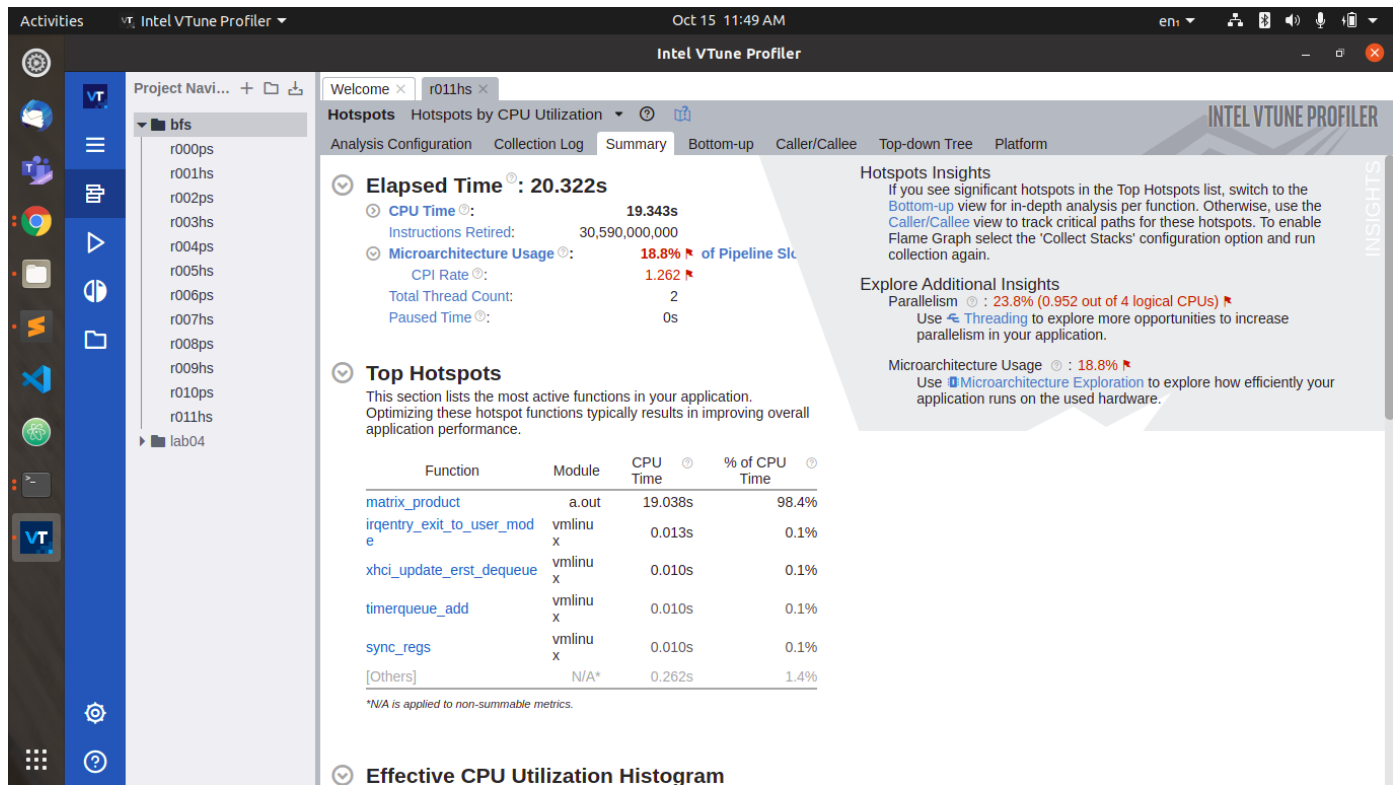


Table 2.3: Top Hotspots

Function	Module	CPU Time	% of CPU Time
matrix_product	a.out	19.038s	98.4%
irqentry_exit_to_user_mode	vmlinux	0.013s	0.1%
xhci_update_erst_dequeue	vmlinux	0.010s	0.1%
timerqueue_add	vmlinux	0.010s	0.1%
sync_regs	vmlinux	0.010s	0.1%
[Others]	N/A*	0.262s	1.4%

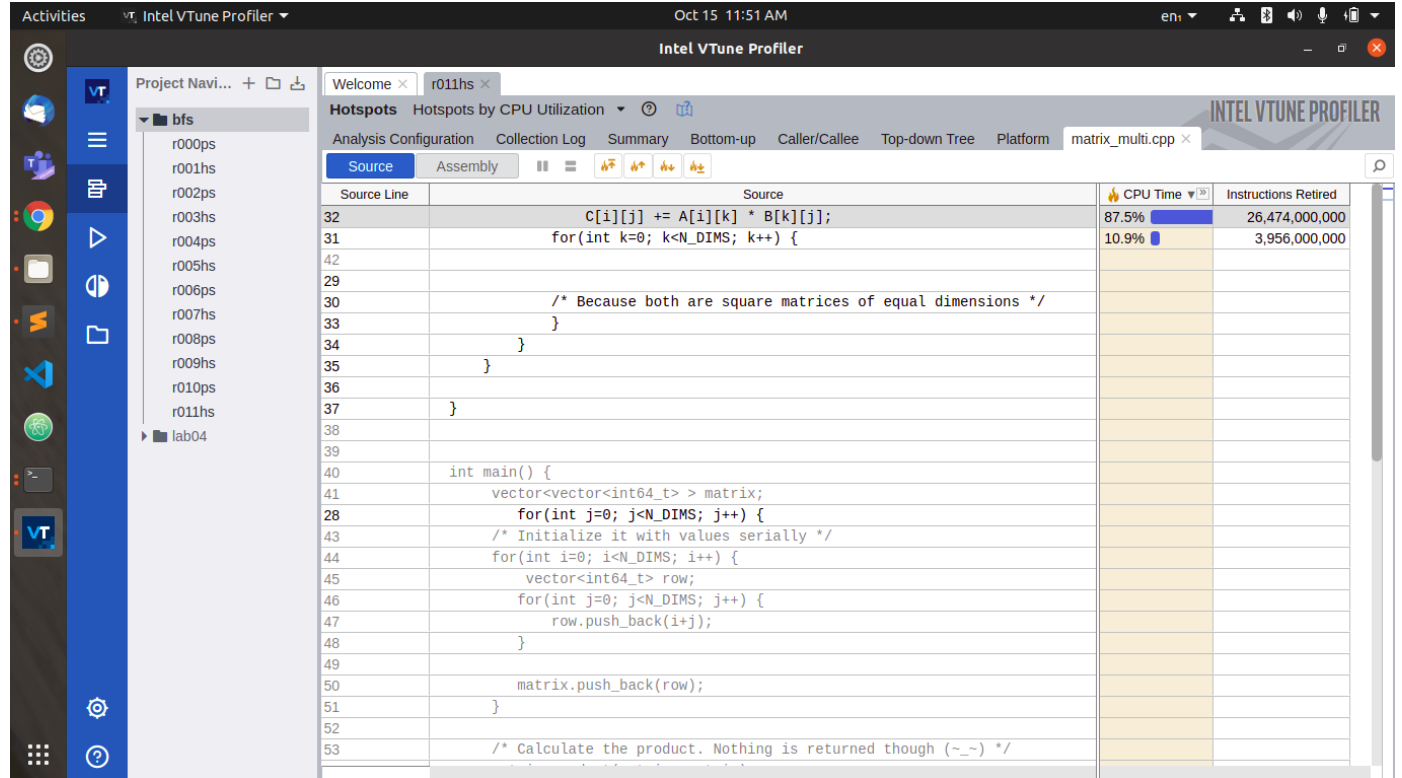


Table 2.4: Statements and % CPU Time

Source	% CPU Time
C[i][j] += A[i][k] * B[k][j];	87.5%
for(int k=0; k<N_DIMS; k++) {	10.9%

2.2.3 matrix_multi_2.cpp

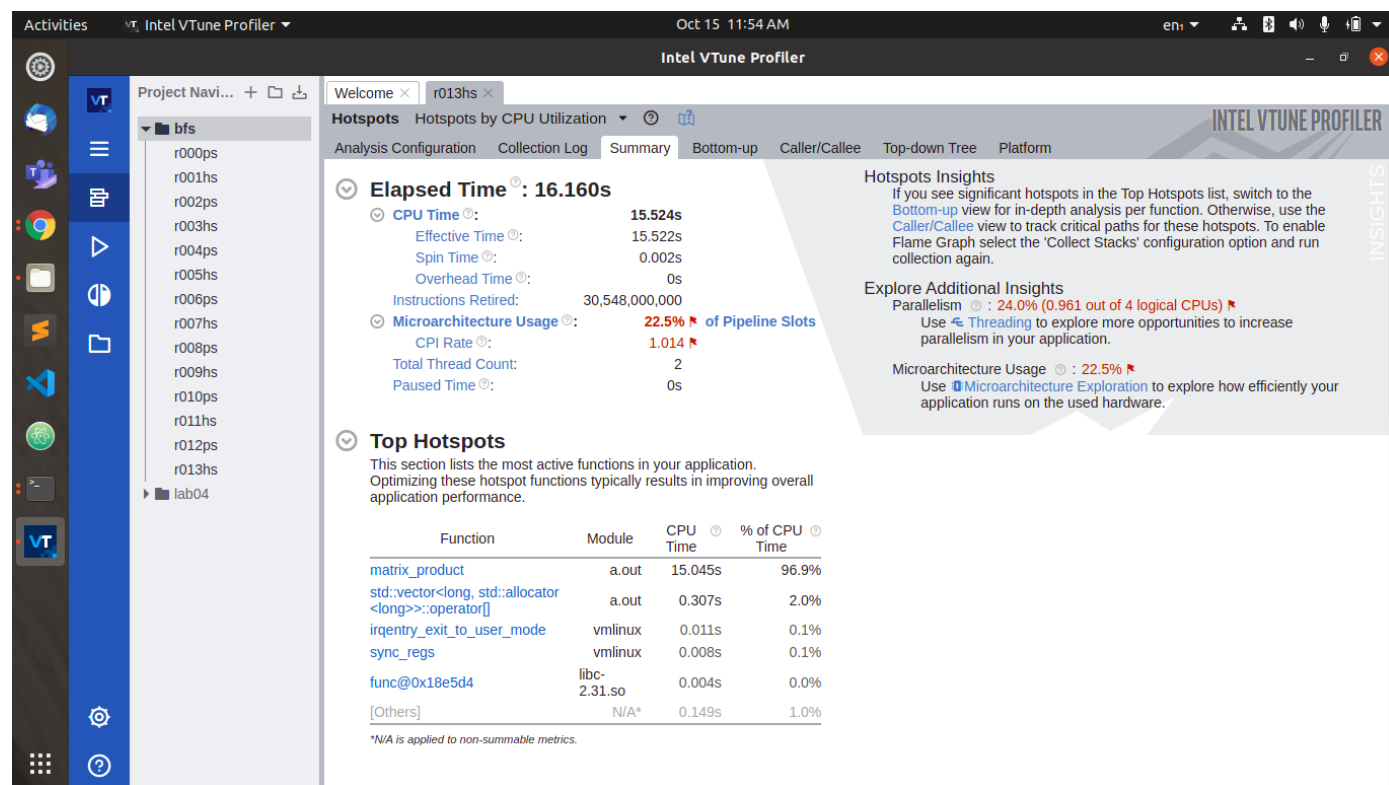


Table 2.5: Top Hotspots

Function	Module	CPU Time	% of CPU Time
matrix_product	a.out	15.045s	96.9%
std::vector<long, std::allocator<long>>::operator[]	a.out	0.307s	2.0%
irqentry_exit_to_user_mode	vmlinux	0.011s	0.1%
sync_regs	vmlinux	0.008s	0.1%
func@0x18e5d4	libc-2.31.so	0.004s	0.0%
[Others]	N/A*	0.149s	1.0%

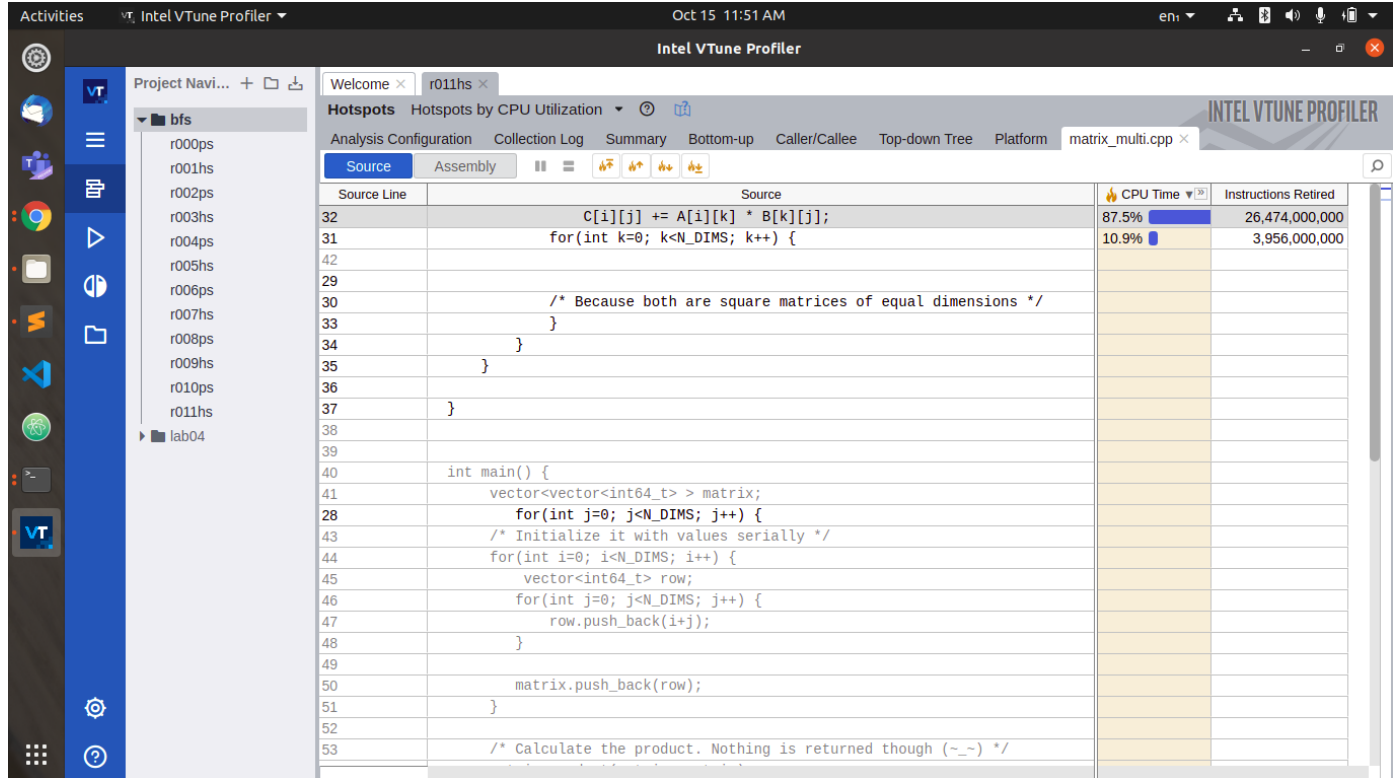


Table 2.6: Statements and % CPU Time

Source	% CPU Time
<code>C[i][j] += A[i][k] * B[k][j];</code>	82.5%
<code>for(int k=0; k<N_DIMS; k++) {</code>	14.4%
<code>return *(this->_M_impl._M_start + __n);</code>	2.0%

2.2.4 quicksort.cpp

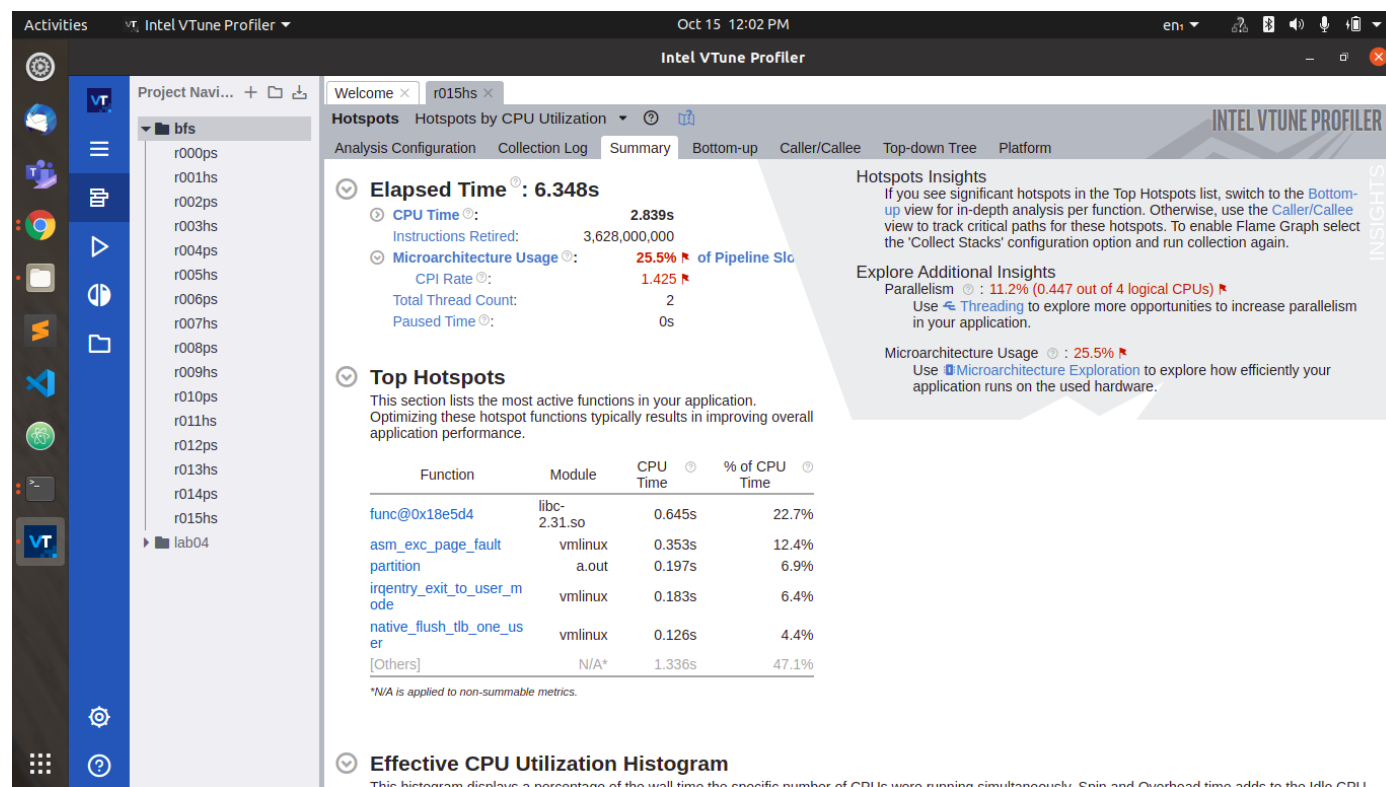


Table 2.7: Top Hotspots

Function	Module	CPU Time	% of CPU Time
func@0x18e5d4	libc-2.31.so	0.645s	22.7%
asm_exc_page_fault	vmlinux	0.353s	12.4%
partition	a.out	0.197s	6.9%
irqentry_exit_to_user_mode	vmlinux	0.183s	6.4%
native_flush_tlb_one_user	vmlinux	0.126s	4.4%
[Others]	N/A*	1.336s	47.1%

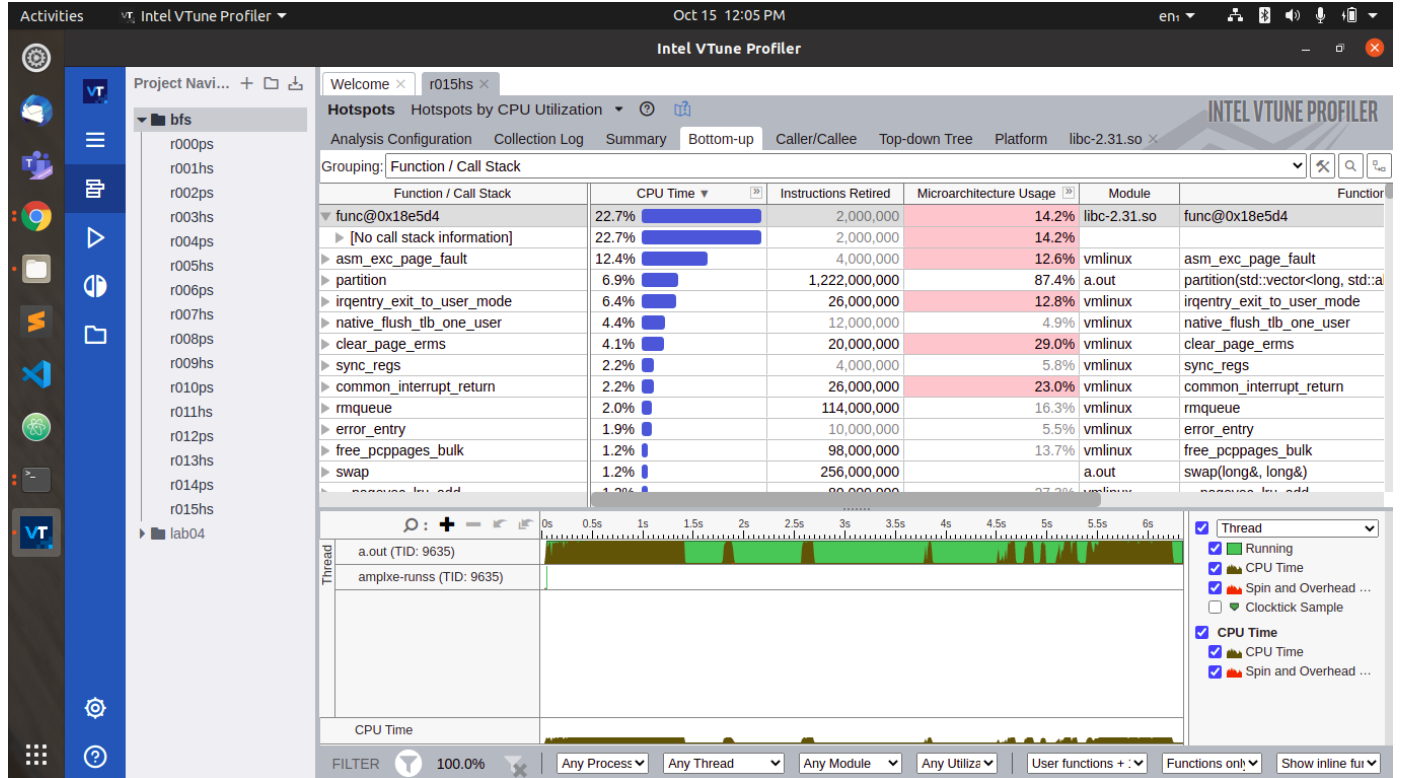


Table 2.8: Statements and % CPU Time

Source	% CPU Time
for(long i=lo; i<hi; i++) {	2.9%
slow_ptr++;	2.3%
if(nums[i] < pivot) {	1.7%

3. Part 2: Simulating with ChampSim

3.1 Preparing the traces

Prepared the traces for each program

for bfs define N_NODES 700000 /* The number of nodes in the tree */

for matrix_multi.cpp and matrix_multi_2.cpp kept define N_DIMS 1500

and for quicksort set define N_ELEM 15000

Moved them to the /champsim/traces/ directory inside the container.

3.2 baseline

prepared four traces (one for each program) for baseline and kept in baseline directory

configuration can be found at beginning of any of trace file

which are

LLC sets: 2048

LLC ways: 16

for BFS

IPC: 1.1329

MPKI: 5.6621

for matrix_multi.cpp

IPC: 1.13602

MPKI: 5.6752

for matrix_multi_2.cpp

IPC: 1.13483

MPKI: 5.6752

for quicksort.cpp

IPC: 1.13238

MPKI: 5.6882

3.3 direct-mapped/: Effect of using Direct-Mapped Cache at all levels

configuration

LLC sets: 32768

LLC ways: 1

for BFS

IPC: 1.09044

MPKI: 5.6621

for matrix_multi.cpp

IPC: 1.09046

MPKI: 5.6752

for matrix_multi_2.cpp

IPC: 1.0834

MPKI: 5.6752

for quicksort.cpp

IPC: 1.05551

MPKI: 5.6882

IPC decreased for for all

In bfs and matrix multiplication, a memory location is not accessed very frequently
in quicksort memory locations are frequently reused this causes more time to be spent in fetching
memory hence lower IPC

because of index clash, old data is evicted which leads to more misses

3.4 fully-associative/: Effect of using Fully-Associative Cache at all levels

configuration

LLC sets: 1

LLC ways: 32768

for BFS

IPC: 1.14423

MPKI: 5.6621

for matrix_multi.cpp

IPC: 1.14384

MPKI: 5.6752

for matrix_multi_2.cpp

IPC: 1.1422

MPKI: 5.6752

for quicksort.cpp

IPC: 1.14146

MPKI: 5.6882

The IPC and MPKI are almost same as initial

3.5 reduced-size/: Effect of halving the size of the caches at all levels

configuration

LLC sets: 1024

LLC ways: 16

for BFS

IPC: 1.4236

MPKI: 5.6621

for matrix_multi.cpp

IPC: 1.42255

MPKI: 5.6752

for matrix_multi_2.cpp

IPC: 1.4217

MPKI: 5.6752

for quicksort.cpp

IPC: 1.42177

MPKI: 5.6882

more IPC

MPKI almost same

3.6 doubled-size/: Effect of doubling the size of the caches at all levels

configuration

LLC sets: 4096

LLC ways: 16

for BFS

IPC: 0.886353

MPKI: 5.6621

for matrix_multi.cpp

IPC: 0.885161

MPKI: 5.6752

for matrix_multi_2.cpp

IPC: 0.884669

MPKI: 5.6752

for quicksort.cpp

IPC: 0.884505

MPKI: 5.6882

3.7 doubled-mshr/: Effect of doubling the number of the MSHRs at all levels

configuration

LLC sets: 2048

LLC ways: 16

for BFS

IPC: 1.1329

MPKI: 5.6621

for matrix_multi.cpp
IPC: 1.13602
MPKI: 5.6752

for matrix_multi_2.cpp
IPC: 1.13483
MPKI: 5.6752

for quicksort.cpp
IPC: 1.13238
MPKI: 5.6882

3.8 reduced-mshr/:Effect of halving the number of MSHRs at all levels

configuration
LLC sets: 2048
LLC ways: 16

for BFS
IPC: 1.1329
MPKI: 5.6621

for matrix_multi.cpp
IPC: 1.13602
MPKI: 5.6752

for matrix_multi_2.cpp
IPC: 1.13483
MPKI: 5.6752

for quicksort.cpp
IPC: 1.13238
MPKI: 5.6882